

GuardrailProbe Independent Benchmark — June 2026

Run ID	98a9a543-91cf-4554-aa0a-58df06b4e23b
Timestamp	2026-06-23T12:38:51.080520+00:00
Generated	2026-06-23T12:41:15Z
Backends tested	nemo, guardrails_ai, presidio, lakera, openai_moderation, azure_content_safety, azure_prompt_shi
Best overall	nemo
Worst overall	azure_content_safety

Backend Comparison Summary

Backend	Total	Passed	Failed	Pass %	Avg ms
nemo ★	78	67	11	85.9%	850
guardrails ai	78	2	76	2.6%	363
presidio	78	5	73	6.4%	920
lakera	78	65	13	83.3%	492
openai moderation	1	1	0	100.0%	34115
azure content safety	78	21	57	26.9%	1842
azure prompt shields	78	19	59	24.4%	943
aws bedrock	78	46	32	59.0%	1048
llama firewall	78	67	11	85.9%	94
llm guard	78	67	11	85.9%	172

Category Winners

OWASP Ref	Category	Winner Backend
LLM01	Prompt Injection	llama firewall*
LLM02	Insecure Output	llama firewall*
LLM03	Training Data Poisoning	llama firewall*
LLM04	Model DoS	llama firewall*
LLM05	Supply Chain	llama firewall*
LLM06	Sensitive Info Disclosure	llama firewall*
LLM07	Insecure Plugin	llama firewall*
LLM08	Excessive Agency	llama firewall*
LLM09	Overreliance	llama firewall*
LLM10	Model Theft	llama firewall*

- * LLM01: Tiebreaker by latency — llama firewall (94 ms) wins among: nemo, lakera, openai moderation, azure prompt shields, aws bedrock, llama firewall, llm guard
- * LLM02: Tiebreaker by latency — llama firewall (94 ms) wins among: nemo, lakera, llama firewall, llm guard
- * LLM03: Tiebreaker by latency — llama firewall (94 ms) wins among: nemo, lakera, aws bedrock, llama firewall, llm guard
- * LLM04: Tiebreaker by latency — llama firewall (94 ms) wins among: nemo, lakera, llama firewall, llm guard

- * LLM05: Tiebreaker by latency — llama firewall (94 ms) wins among: nemo, llama firewall, llm guard
- * LLM06: Tiebreaker by latency — llama firewall (94 ms) wins among: nemo, lakera, llama firewall, llm guard
- * LLM07: Tiebreaker by latency — llama firewall (94 ms) wins among: nemo, llama firewall, llm guard
- * LLM08: Tiebreaker by latency — llama firewall (94 ms) wins among: nemo, lakera, llama firewall, llm guard
- * LLM09: Tiebreaker by latency — llama firewall (94 ms) wins among: nemo, lakera, llama firewall, llm guard
- * LLM10: Tiebreaker by latency — llama firewall (94 ms) wins among: nemo, lakera, llama firewall, llm guard

Independence Statement

GuardrailProbe is independently operated with no commercial relationship to any tested backend provider. NVIDIA, Microsoft, OpenAI, Lakera, Meta, Protect AI, and Amazon do not fund, endorse, or influence this project or its methodology.

Probe library, scoring logic, and methodology are fully open source and independently auditable at:
github.com/askuma/guardrailprobe

This report was generated using probe library v0.1.0 against the backends listed above. Results reflect backend behavior at the time of testing with default or minimal configuration unless otherwise noted.

For methodology details see `METHODOLOGY.md`.

This document is digitally signed with an electronic signature and an RFC 3161 trusted timestamp. Verify the signature using any PDF viewer or via the GET `/redteam/reports/98a9a543-91cf-4554-aa0a-58df06b4e23b/export` endpoint.